



Music & Speech Detection in Radio Broadcasts

Table of Content

*Things to Note: Due to weak computer, not all files provided was used, a portion was randomly chosen and used to train classifier



Diagrams created based on data



Discussion on Features, Outliers and Distributions of Result



Most accurate algorithm



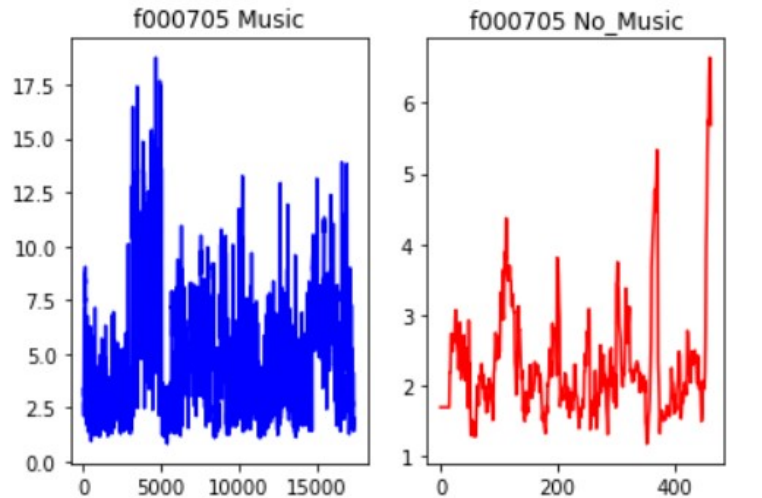
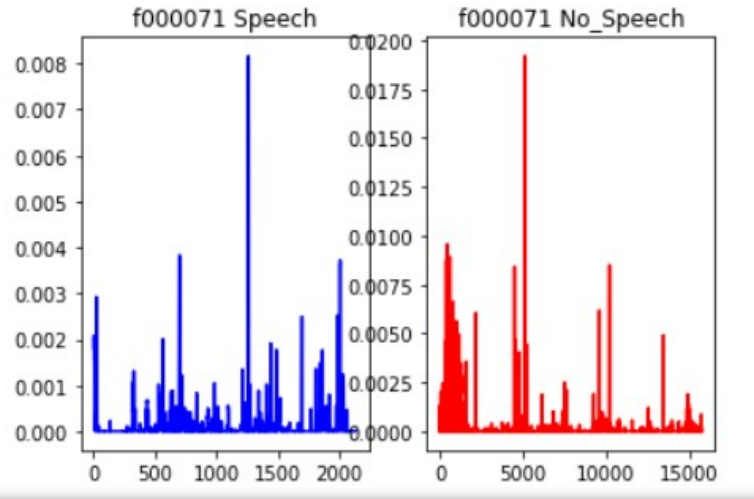
Observation on overfitting behavior



Variation in Classification accuracy based on relevant parameters

Informative & Uninformative Features

- When there is music, there is generally a higher frequency range than when it has no music.
 - F000657 No_music, frequency limited to between -0.10 and 0.20
- Number of rows detected classifying as speech or no music is generally shorter than “music” or “no speech”.
- Opposites groups (music, no speech) vs (no music, speech)



A brief look on correlation

Due to less
features on
Speech than
Music, correlation
is generally lower

Correlation done
on features,
showing that
correlation
between columns
nearer each other
are more similar
to each other
than those
further apart



Pearson Correlation Music

f000453	f000101	-0.898801
f000101	f000453	-0.898801
f000064	f000416	-0.896595
f000416	f000064	-0.896595
	f000063	-0.886645
f000063	f000416	-0.886645
f000017	f000370	-0.877621
f000370	f000017	-0.877621
f000083	f000435	-0.877394
f000435	f000083	-0.877394
f000646	f000294	-0.870709
f000294	f000646	-0.870709
f000425	f000073	-0.865146
f000073	f000425	-0.865146
f000404	f000052	-0.857129
f000052	f000404	-0.857129
f000404	f000051	-0.856838
f000051	f000404	-0.856838
f000622	f000270	-0.839810
f000270	f000622	-0.839810

f000076	f000077	0.908792
f000077	f000076	0.908792
f000026	f000027	0.912269
f000027	f000026	0.912269
f000082	f000081	0.917242
f000081	f000082	0.917242
f000099	f000100	0.918333
f000100	f000099	0.918333
f000059	f000058	0.922824
f000058	f000059	0.922824
f000087	f000088	0.925209
f000088	f000087	0.925209
f000041	f000042	0.931993
f000042	f000041	0.931993
f000052	f000051	0.943243
f000051	f000052	0.943243
f000044	f000045	0.957130
f000045	f000044	0.957130
f000627	f000276	0.961130
f000276	f000627	0.961130

Pearson Correlation Speech

f000065	f000077	-0.083020
f000077	f000065	-0.083020
f000078	f000071	-0.071524
f000071	f000078	-0.071524
f000077	f000068	-0.063018
f000068	f000077	-0.063018
f000077	f000064	-0.056686
f000064	f000077	-0.056686
f000077	f000071	-0.055927
f000071	f000077	-0.055927
f000077	f000072	-0.055296
f000072	f000077	-0.055296
	f000078	-0.054866
f000078	f000072	-0.054866
f000061	f000077	-0.043882
f000077	f000061	-0.043882
f000072	f000049	-0.041894
f000049	f000072	-0.041894
f000071	f000049	-0.041649
f000049	f000071	-0.041649

f000039	f000038	0.696343
f000038	f000039	0.696343
f000023	f000022	0.698257
f000022	f000023	0.698257
f000034	f000035	0.699175
f000035	f000034	0.699175
f000018	f000017	0.702901
f000017	f000018	0.702901
f000008	f000007	0.703285
f000007	f000008	0.703285
f000008	f000009	0.703443
f000009	f000008	0.703443
f000033	f000034	0.707122
f000034	f000033	0.707122
f000026	f000025	0.710794
f000025	f000026	0.710794
f000017	f000016	0.711793
f000016	f000017	0.711793
f000042	f000043	0.716122
f000043	f000042	0.716122

Outliers found in Data

Outliers found using statistical method, with ROC as line of best fit separating

Music

KNN

On Training Data:

KNN ROC:0.1353, precision @ rank n:0.0031

Outliers 321

Inliers 12204

On Test Data:

KNN ROC:0.1474, precision @ rank n:0.0071

Outliers 141

Inliers 5227

Speech

KNN

On Training Data:

KNN ROC:0.8818, precision @ rank n:0.5152

Outliers 1549

Inliers 10976

On Test Data:

KNN ROC:0.8869, precision @ rank n:0.4805

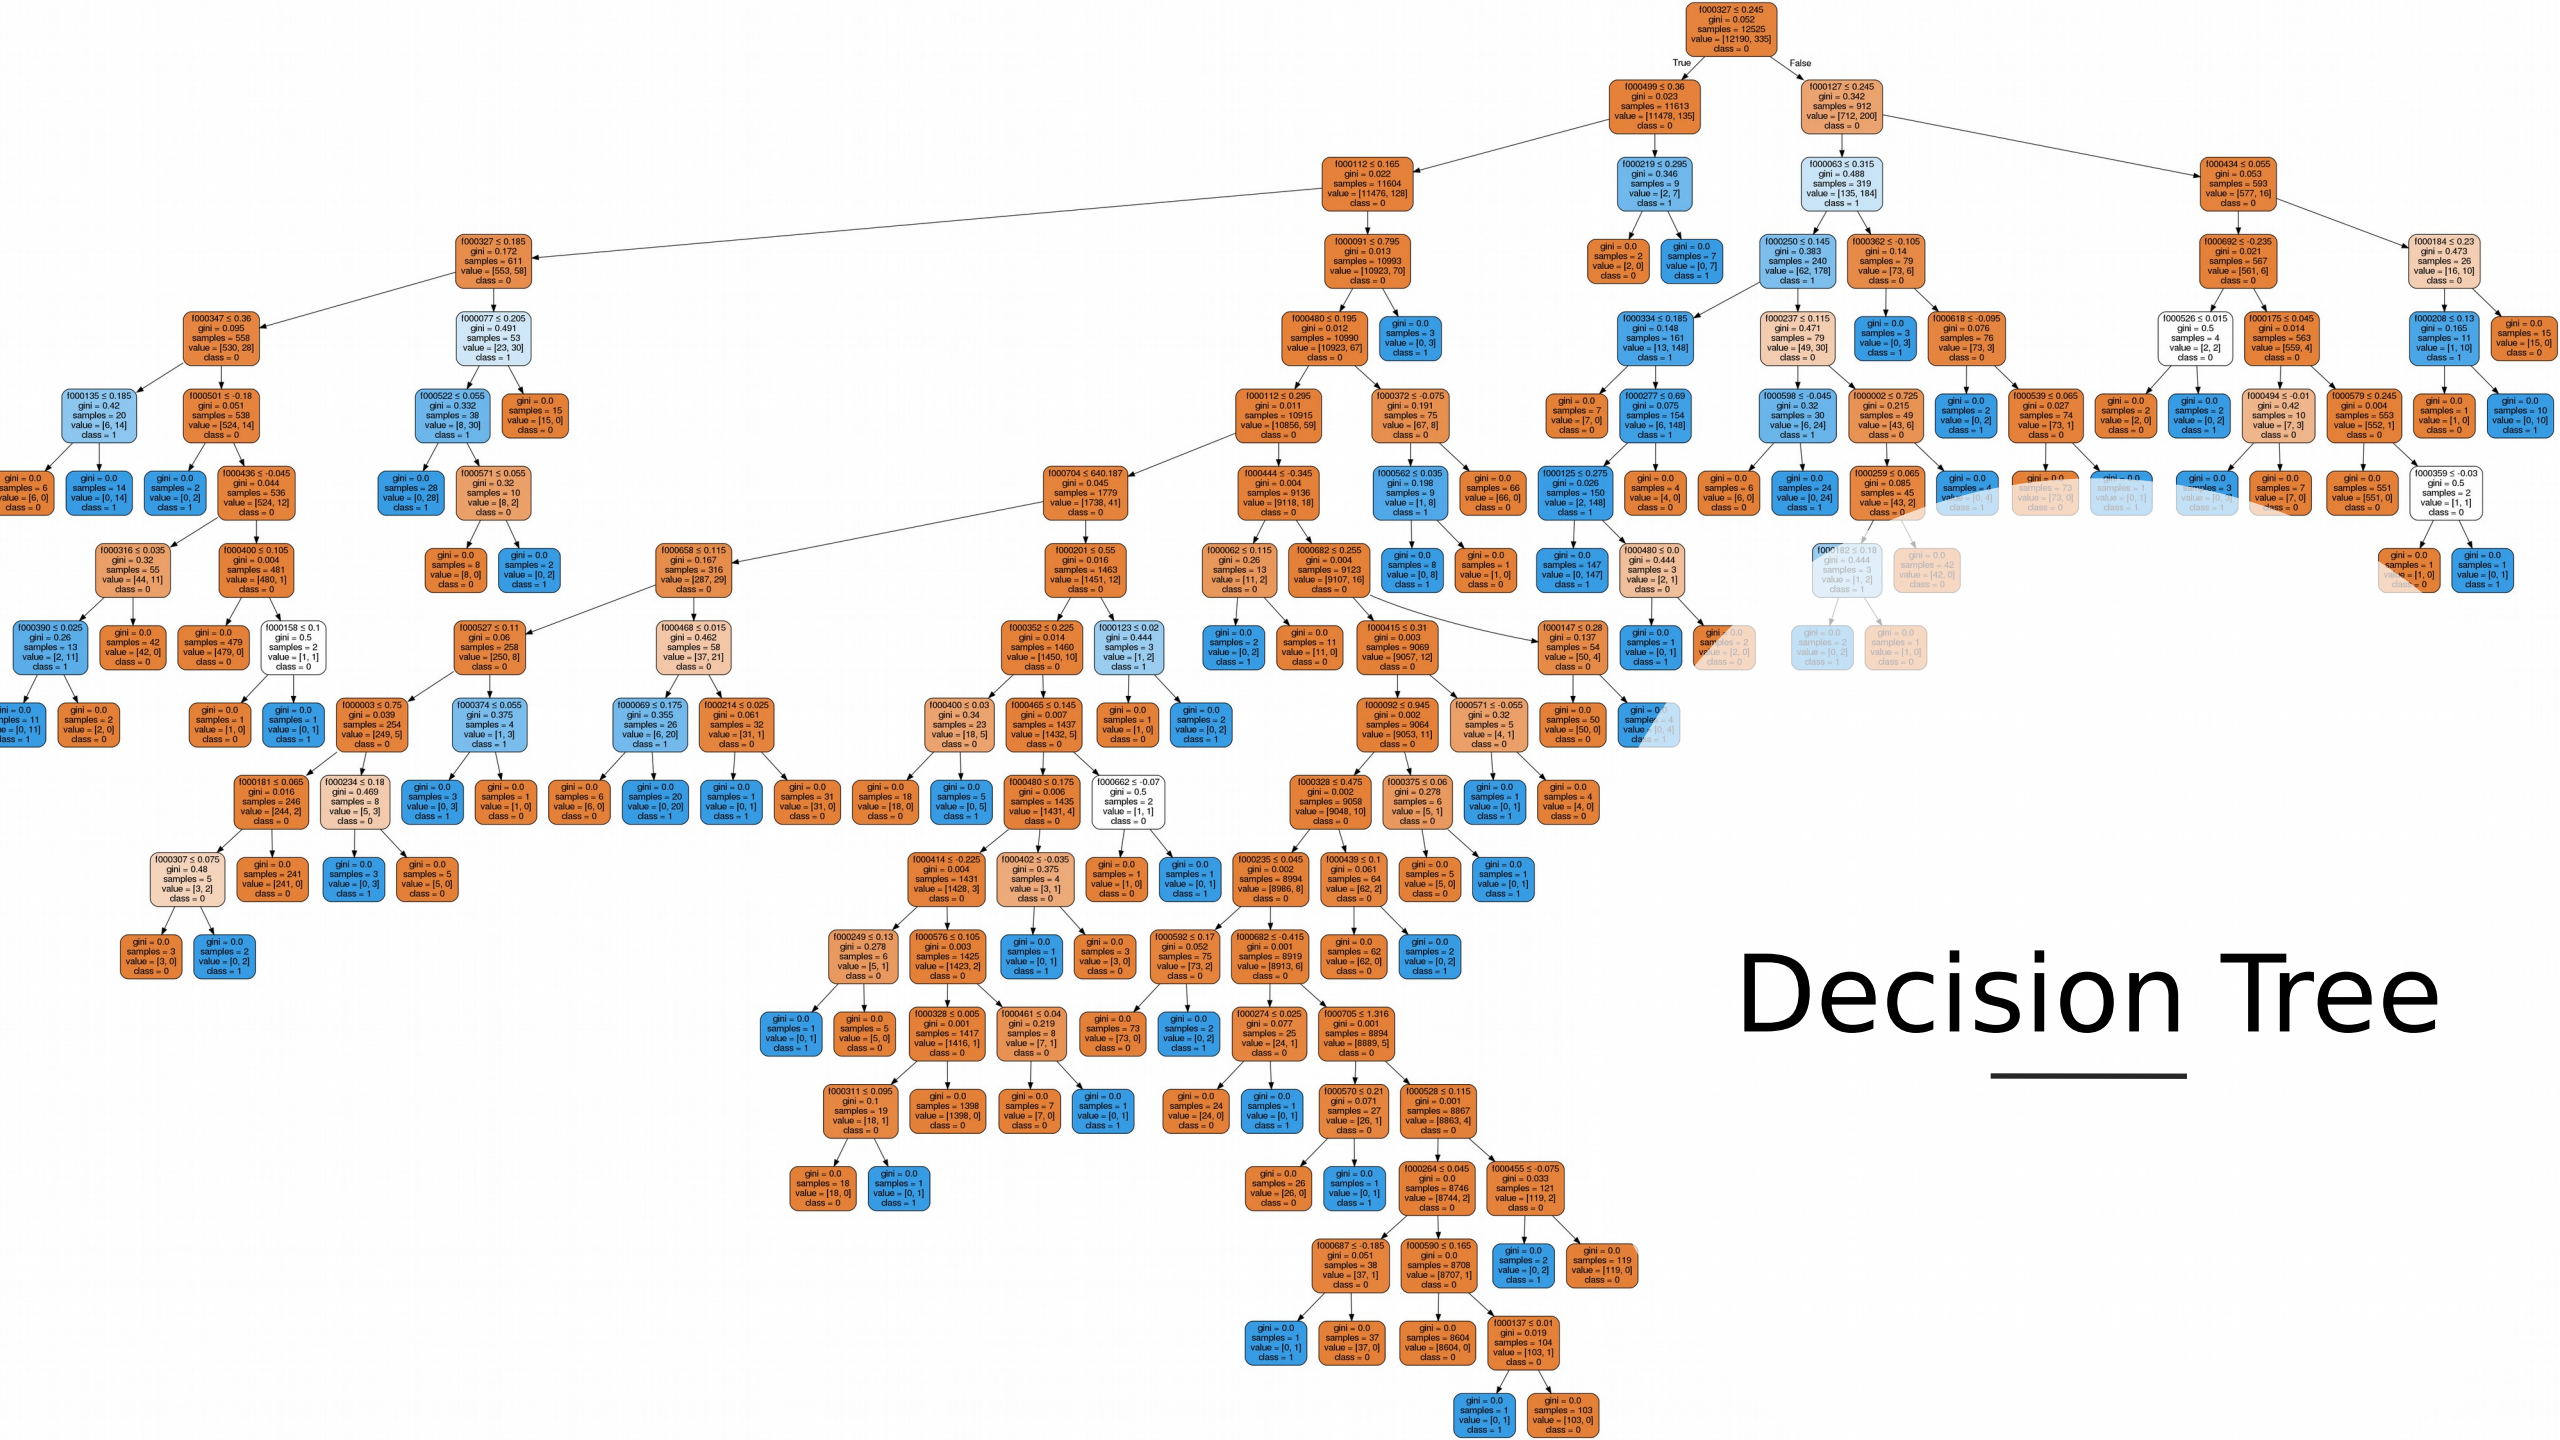
Outliers 591

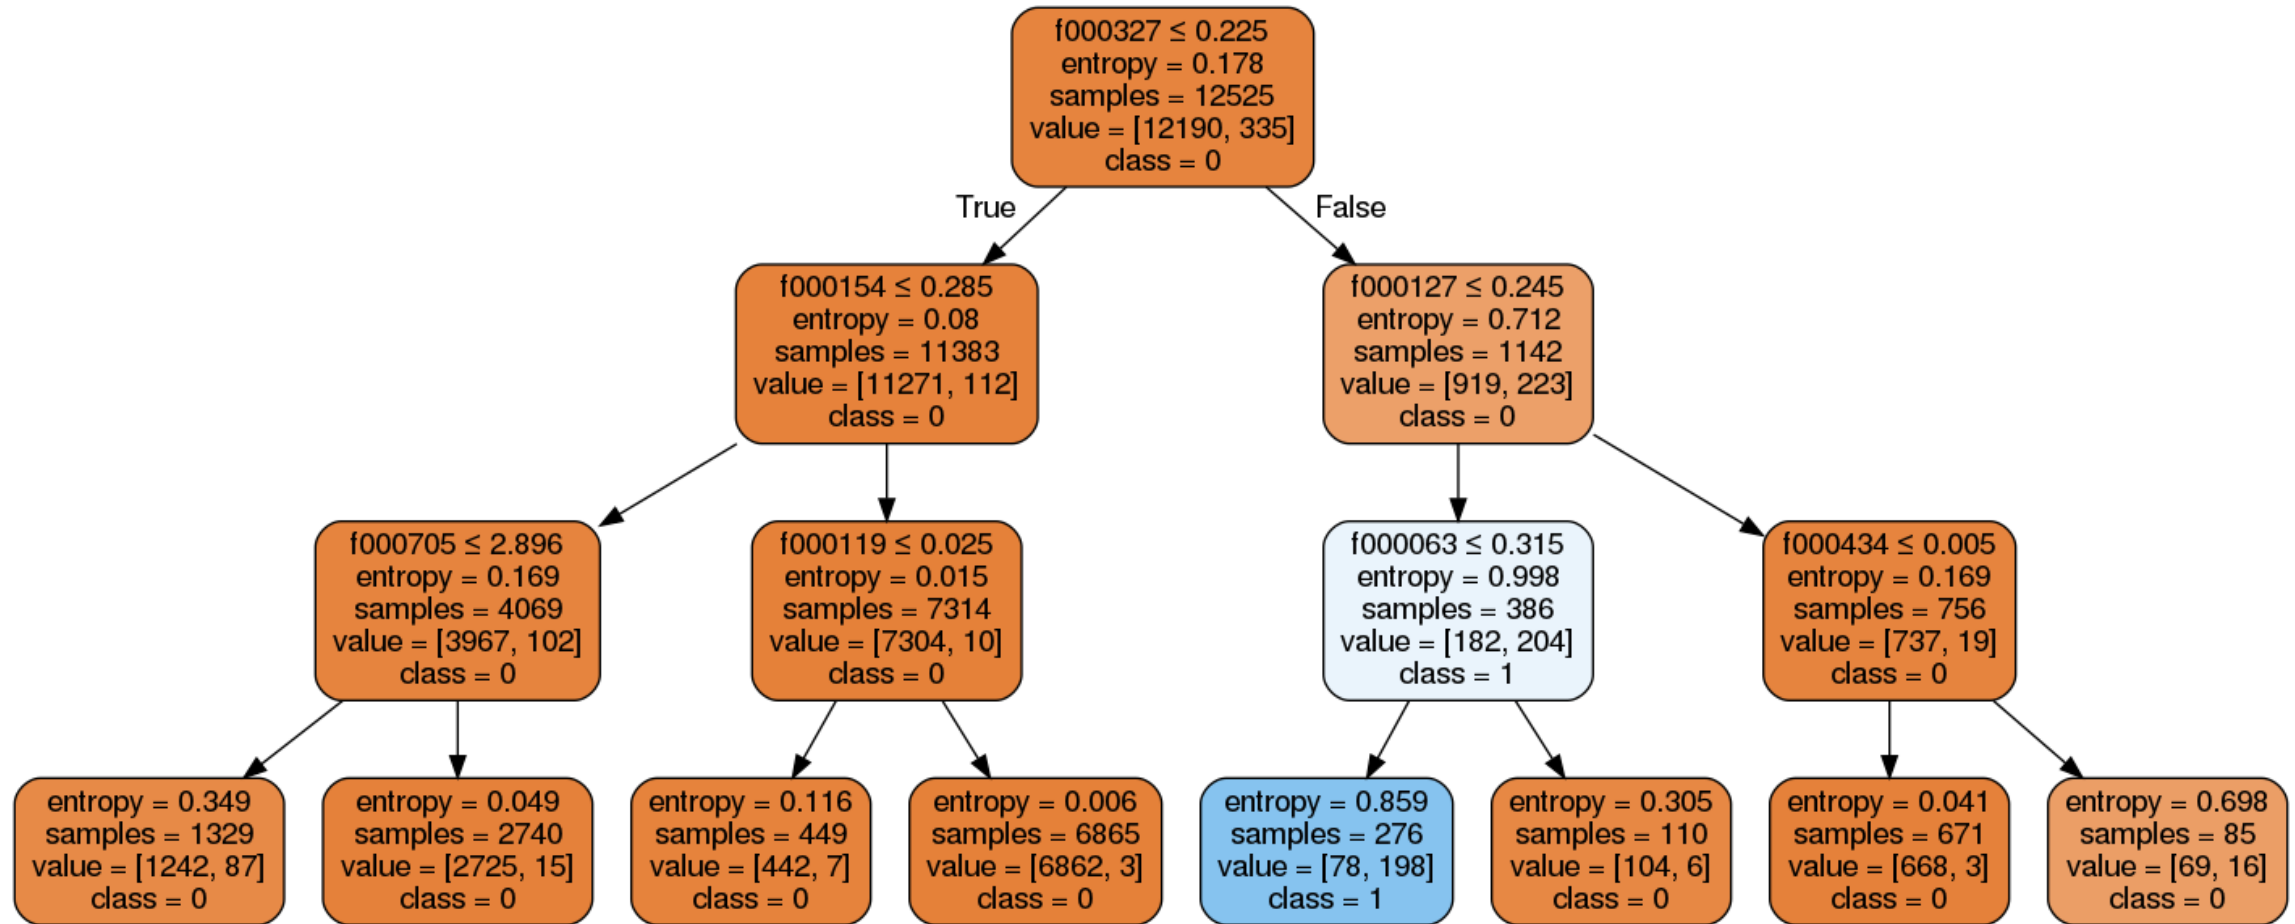
Inliers 4777

Most accurate algorithm

```
models = [  
    ('LR', LogisticRegression()),  
    ('NB', GaussianNB()),  
    ('SVM', SVC()),  
    ('KNN', KNeighborsClassifier()),  
    ('DT', DecisionTreeClassifier()),  
    ('RF', RandomForestClassifier()),  
]
```

- Depending on either music or speech, but generally, logistic regression or random forest
- This is further validated with testing using feature selection, where logistic regression or random forest remains most accurate





Simple fitting and scores

Music

```
for name, model in models:  
    clf = model  
    clf.fit(X_train, y_train)  
    accuracy = clf.score(X_test, y_test)  
    print(name, accuracy)
```

LR 0.9953427719821163
NB 0.9329359165424739
SVM 0.9726154992548435
KNN 0.9673994038748137
DT 0.9897540983606558
RF 0.9927347242921013

Speech

```
for name, model in models:  
    clf = model  
    clf.fit(X_train, y_train)  
    accuracy = clf.score(X_test, y_test)  
    print(name, accuracy)
```

LR 0.9210134128166915
NB 0.8444485842026825
SVM 0.8764903129657228
KNN 0.897354694485842
DT 0.9072280178837556
RF 0.937220566318927

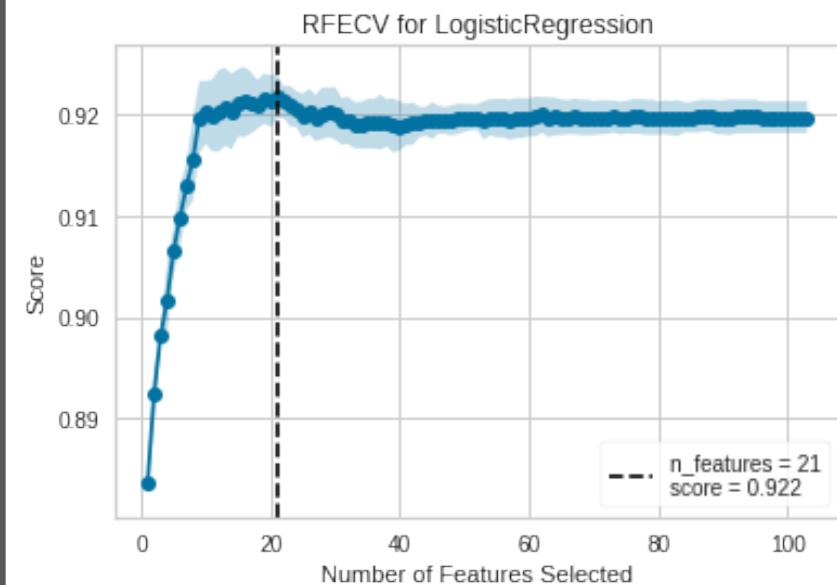

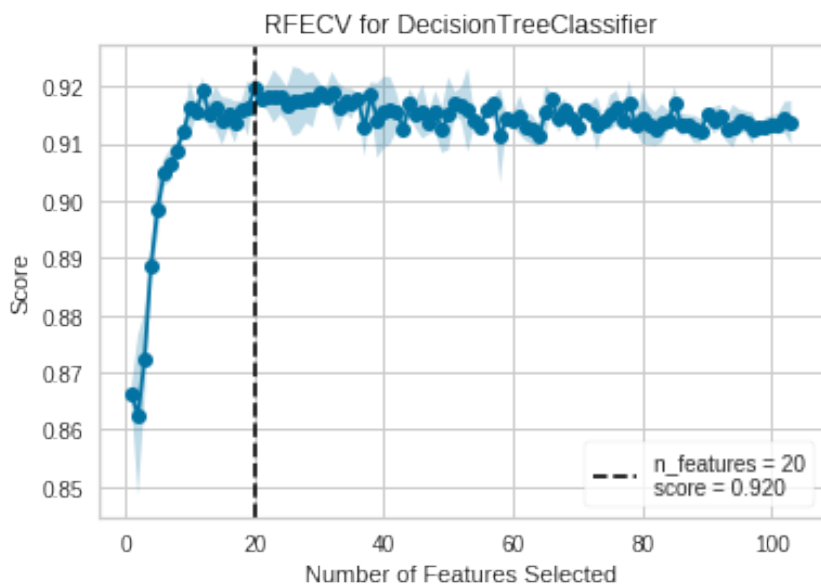
Cross validation and Average of accuracy

Music

```
LR [0.96284916 0.95110366 0.98211291 0.96394634 0.99357183] Accuracy: 0.97 (+/- 0.03)
NB [0.91312849 0.92204526 0.98798211 0.87562884 0.96646171] Accuracy: 0.93 (+/- 0.08)
SVM [0.97402235 0.97401509 0.97428731 0.97428731 0.97428731] Accuracy: 0.97 (+/- 0.00)
KNN [0.96256983 0.95082425 0.97540525 0.93236445 0.97931805] Accuracy: 0.96 (+/- 0.03)
DT [0.93463687 0.92316289 0.93935159 0.95248742 0.98071548] Accuracy: 0.95 (+/- 0.04)
RF [0.98100559 0.95920648 0.97959754 0.97540525 0.980436 ] Accuracy: 0.98 (+/- 0.02)
```

Speech

```
LR [0.90611903 0.89857502 0.92875105 0.91196199 0.9175517 ] Accuracy: 0.91 (+/- 0.02)
NB [0.84492875 0.79519419 0.86784018 0.85997764 0.84935718] Accuracy: 0.84 (+/- 0.05)
SVM [0.88041352 0.88041352 0.88041352 0.8803801 0.8803801 ] Accuracy: 0.88 (+/- 0.00)
KNN [0.89354568 0.88851635 0.89801621 0.90804919 0.88177753] Accuracy: 0.89 (+/- 0.02)
DT [0.8577815 0.8611344 0.89606035 0.87311347 0.8627725 ] Accuracy: 0.87 (+/- 0.03)
RF [0.8977368 0.89801621 0.92763342 0.92733371 0.91224148] Accuracy: 0.91 (+/- 0.03)
```

Recursive Feature Elimination Cross Validation – Speech

Feature importance Random forest - music

f000328	0.018801
f000327	0.017725
f000350	0.012909
f000112	0.011542
f000329	0.011431
f000113	0.010778
f000331	0.010410
f000345	0.010282
f000099	0.009811
f000154	0.008757
f000351	0.008111
f000077	0.008070
f000181	0.008041
f000699	0.007767
f000705	0.007697
f000452	0.007394
f000162	0.006931
f000128	0.006702

Overfitting KNeighbors - speech

Simple fitting and prediction

When trying to best feature, initial results seem to imply that K-neighbors was most accurate algorithm, however, further observations using cross validating reveals it was overfitted, lowering it's score.

```
algorithm: LR , accuracy: 0.8841870824053452 , number of features 2
algorithm: RF , accuracy: 0.884558277654046 , number of features 2
algorithm: LR , accuracy: 0.8971789161098738 , number of features 3
algorithm: NB , accuracy: 0.897735708982925 , number of features 3
algorithm: SVM , accuracy: 0.9007052709725315 , number of features 3
algorithm: LR , accuracy: 0.9072011878247959 , number of features 4
algorithm: NB , accuracy: 0.9073867854491463 , number of features 4
algorithm: KNN , accuracy: 0.9090571640683 , number of features 5
algorithm: LR , accuracy: 0.9142538975501113 , number of features 6
algorithm: KNN , accuracy: 0.9192650334075724 , number of features 6
algorithm: LR , accuracy: 0.9218634001484781 , number of features 7
algorithm: KNN , accuracy: 0.9304008908685969 , number of features 7
algorithm: KNN , accuracy: 0.9318856718634001 , number of features 8
algorithm: KNN , accuracy: 0.9367112100965108 , number of features 9
algorithm: RF , accuracy: 0.9426503340757239 , number of features 9
algorithm: KNN , accuracy: 0.9452487008166296 , number of features 13
algorithm: KNN , accuracy: 0.950445434298441 , number of features 16
algorithm: KNN , accuracy: 0.9523014105419451 , number of features 17
algorithm: KNN , accuracy: 0.955456570155902 , number of features 19
algorithm: RF , accuracy: 0.9556421677802525 , number of features 19
algorithm: KNN , accuracy: 0.9615812917594655 , number of features 20
```

LR	[0.93290646 0.92010022 0.86358575 0.82126949 0.80874165]	Accuracy: 0.87 (+/- 0.10)	, number of features: 1
LR	[0.93012249 0.92622494 0.89142539 0.85467706 0.81375278]	Accuracy: 0.88 (+/- 0.09)	, number of features: 2
DT	[0.93402004 0.93374165 0.88864143 0.85662584 0.81904232]	Accuracy: 0.89 (+/- 0.09)	, number of features: 2
LR	[0.94961024 0.93541203 0.88864143 0.86414254 0.83240535]	Accuracy: 0.89 (+/- 0.09)	, number of features: 3
SVM	[0.9535078 0.94209354 0.8905902 0.86108018 0.83101336]	Accuracy: 0.90 (+/- 0.09)	, number of features: 3
LR	[0.94654788 0.94821826 0.9064588 0.87722717 0.8516147]	Accuracy: 0.91 (+/- 0.08)	, number of features: 4
LR	[0.94738307 0.95267261 0.90896437 0.875 0.85885301]	Accuracy: 0.91 (+/- 0.08)	, number of features: 5
LR	[0.95128062 0.95517817 0.90673719 0.87193764 0.86831849]	Accuracy: 0.91 (+/- 0.07)	, number of features: 6
LR	[0.95851893 0.96603563 0.9064588 0.87917595 0.87889755]	Accuracy: 0.92 (+/- 0.08)	, number of features: 7
LR	[0.95629176 0.96798441 0.90979955 0.88585746 0.88223831]	Accuracy: 0.92 (+/- 0.07)	, number of features: 8
LR	[0.95657016 0.9685412 0.91035635 0.88613586 0.88363029]	Accuracy: 0.92 (+/- 0.07)	, number of features: 9
RF	[0.95211581 0.96380846 0.90924276 0.90367483 0.87945434]	Accuracy: 0.92 (+/- 0.06)	, number of features: 9
RF	[0.95072383 0.95824053 0.90868597 0.91146993 0.88446548]	Accuracy: 0.92 (+/- 0.06)	, number of features: 10
RF	[0.95657016 0.96046771 0.90979955 0.90311804 0.88585746]	Accuracy: 0.92 (+/- 0.06)	, number of features: 13
RF	[0.95768374 0.96269488 0.91063474 0.905902 0.8908686]	Accuracy: 0.93 (+/- 0.06)	, number of features: 14
RF	[0.95517817 0.9607461 0.91035635 0.91119154 0.89643653]	Accuracy: 0.93 (+/- 0.05)	, number of features: 15
RF	[0.95824053 0.96018931 0.91063474 0.9136971 0.90005568]	Accuracy: 0.93 (+/- 0.05)	, number of features: 18
RF	[0.95935412 0.95740535 0.91731626 0.91341871 0.90423163]	Accuracy: 0.93 (+/- 0.05)	, number of features: 20

With a Cross validation K Neighbors does not have the best accuracy

Random Forest:

Best param: {'n_estimators': 1400, 'min_samples_split': 2, 'min_samples_leaf': 1, 'max_features': 'auto', 'max_depth': 40, 'bootstrap': False}

Best estimator RandomForestClassifier(bootstrap=False, class_weight=None, criterion='gini',
max_depth=40, max_features='auto', max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=1400, n_jobs=1,
oob_score=False, random_state=None, verbose=0,
warm_start=False)

Best score 0.9524554735822341

Variations in
classification accuracy
based on parameters

- Through hyperparameter tuning, accuracy of all algorithms are improved.
 - Some notable improvements include Random forest, which saw marked increase of up to 5%

Hyperparameter Tuning - Speech

SVC, K-Neighbors, Logistic Regression, Decision Tree

```
Score 0.9975413500223513
Best param: {'C': 10, 'kernel': 'linear'}
Best estimator SVC(C=10, cache_size=200, class_weight=None, coef0=0.0,
  decision_function_shape='ovr', degree=3, gamma='auto', kernel='linear',
  max_iter=-1, probability=False, random_state=None, shrinking=True,
  tol=0.001, verbose=False)
Best score 0.9976153215589836
```

```
Score 0.9074653553866786
Best param: {'metric': 'minkowski', 'n_neighbors': 3, 'weights': 'uniform'}
Best estimator KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
  metric_params=None, n_jobs=1, n_neighbors=3, p=2,
  weights='uniform')
Best score 0.9023027051196065
```

```
Score 0.934957532409477
Best param: {'C': 100, 'penalty': 'l1', 'solver': 'liblinear'}
Best estimator LogisticRegression(C=100, class_weight=None, dual=False, fit_intercept=True,
  intercept_scaling=1, max_iter=100, multi_class='warn',
  n_jobs=None, penalty='l1', random_state=None, solver='liblinear',
  tol=0.0001, verbose=0, warm_start=False)
Best score 0.9326328340412847
```

```
Score 0.9320518551631649
Best param: {'criterion': 'entropy', 'max_depth': 10000, 'min_samples_leaf': 20}
Best estimator DecisionTreeClassifier(class_weight=None, criterion='entropy',
  max_depth=10000, max_features=None, max_leaf_nodes=None,
  min_impurity_decrease=0.0, min_impurity_split=None,
  min_samples_leaf=20, min_samples_split=2,
  min_weight_fraction_leaf=0.0, presort=False, random_state=None,
  splitter='best')
Best score 0.923690289887473
```


Hyperparameter Tuning - Music

SVC, K-Neighbors, Logistic Regression, Decision Tree, Random Forest

Score 0.9975413500223513

Best param: {'C': 10, 'kernel': 'linear'}

Best estimator SVC(C=10, cache_size=200, class_weight=None, coef0=0.0, decision_function_shape='ovr', degree=3, gamma='auto', kernel='linear', max_iter=-1, probability=False, random_state=None, shrinking=True, tol=0.001, verbose=False)

Best score 0.9976153215589836

Score 0.9984354045596782

Best param: {'C': 10000, 'penalty': 'l1', 'solver': 'liblinear'}

Best estimator LogisticRegression(C=10000, class_weight=None, dual=False, fit_intercept=True, intercept_scaling=1, max_iter=100, multi_class='warn', n_jobs=None, penalty='l1', random_state=None, solver='liblinear', tol=0.0001, verbose=0, warm_start=False)

Best param: {'n_estimators': 1000, 'min_samples_split': 2, 'min_samples_leaf': 1, 'max_features': 'auto', 'max_depth': 50, 'bootstrap': False}

Best estimator RandomForestClassifier(bootstrap=False, class_weight=None, criterion='gini', max_depth=50, max_features='auto', max_leaf_nodes=None, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=1, min_samples_split=2, min_weight_fraction_leaf=0.0, n_estimators=1000, n_jobs=1, oob_score=False, random_state=None, verbose=0, warm_start=False)

Best score 0.9948580371115583

Score 0.9843540455967814

Best param: {'metric': 'manhattan', 'n_neighbors': 3, 'weights': 'distance'}

Best estimator KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='manhattan', metric_params=None, n_jobs=1, n_neighbors=3, p=2, weights='distance')

Score 0.9912829682610639

Best param: {'criterion': 'entropy', 'max_depth': 100, 'min_samples_leaf': 1}


Best estimator DecisionTreeClassifier(class_weight=None, criterion='entropy', max_depth=100, max_features=None, max_leaf_nodes=None, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=1, min_samples_split=2, min_weight_fraction_leaf=0.0, presort=False, random_state=None, splitter='best')

Best score 0.9914300618525971

Calibration of Classifiers

- A side exercise done for the practical was to compare the calibrations of the classifiers.
- As can be observed, logistic regression is the closest to the perfect calibration line. Might be due to logistic regression directly optimizing log-loss, whilst other classifiers gave more biased probabilities

Music



Speech Calibration Plots

